

What is Amazon Elastic Load Balancer (ELB)

Amazon ELB allows you to make your applications highly available by using health checks and distributing traffic across a number of instances.

Consider that you have a WordPress blog which is running on a single t2-micro EC2 instance.

Now you publish an article, it goes viral and your site gets hundreds of thousands of requests. Since you are using a single t2-micro, your website will probably crash.


So, what can you do to avoid this?

You may decide to launch a larger instance like an m5-large in place of t2-micro. This is called vertical scaling when you replace an instance with a more powerful instance.

But vertical scaling isn't always economical.

Another approach can be to use a bunch of smaller instances like t2-micros and distribute the website traffic between them. And Elastic Load Balancer allows you to do just that.

It distributes incoming application or network traffic across multiple targets, such as Amazon EC2 instances, containers, and IP addresses, in multiple Availability Zones.



Credits —aws.amazon.com

It uses health checks to detect which instances are healthy and directs traffic only across those instances.

Types of Elastic Load Balancers

There are three types of load balancers available. You can use the one

that best fits your use case.

1. Classic Load Balancer (CLB)

This is the previous generation load balancer that was used for EC2-classic instances.

It operates on both the request level and the connection level. But it doesn't support features like host-based routing or path-based routing.

Once configured, it distributes the load across all the registered instances regardless of what is present on the servers.

Hence, it can only be used to distribute traffic to a single URL.

2. Application Load Balancer (ALB)

This load balancer is specially designed for web applications with HTTP and HTTPS traffic.

There is a networking model called the OSI Model (Open Systems Interconnection) that is used to explain how computer networks work. This model has 7 layers and the top layer is the Application Layer.

This load balancer works at this Application Layer, hence the name.

It also provides advanced routing features such as host-based and path-based routing and also works with containers and microservices.

Host-based Routing

Suppose you have two websites **medium.com** and **admin.medium.com**. Each website is hosted on two EC2 instances for high availability and you want to distribute the incoming web traffic between them.

If you were using the CLB you would have to create two load balancers, one for each website.

But you can do the same thing using a single ALB!

Hence you will be saving money as you will only be paying for a single ALB instead of two CLBs.

Path-based Routing

Suppose the website of your company is **payzello.com** and the company's blog is hosted on **payzello.com/blog**. The operations team has decided to host the main website and the blog on different instances.

Using ALB you can route traffic based on the path of the requested URL. So again a single ALB is enough to handle this for you.

3. Network Load Balancer (NLB)

This load balancer operates at the Network layer of the OSI model, hence the name.

Suppose your company's website is running on four m4-xlarge instances and you are using an ALB to distribute the traffic among them.

Now your company launched a new product today which got viral and your website starts to get millions of requests per second.

In this case, the ALB may not be able to handle the sudden spike in traffic.

This is where the NLB really shines. It has the capability to handle a sudden spike in traffic since it works at the connection level.

It also provides support for static IPs.

I hope by this time you have got a rough idea about load balancers. Now, enough talking, let's go practical.

Creating an Application Load Balancer

We will handle a case of path-based routing. We will be handling two paths here, "/" and "/blog".

We will launch two instances, one for handling each path. Let's get started!

1. Launch two EC2 instances

To learn how to launch an EC2 instance, you can read my article on [Launching an Amazon EC2 instance](#).

When launching, give a Name tag to your instances.

For the first instance, give a tag with **Name** as key and **Main** as the value. For the second instance, give a tag with **Name** as key and **Blog** as the value. This will help us in distinguishing between them.

After launching the two instances, your dashboard should look like this.

Filter by tags and attributes or search by keyword									
<input type="checkbox"/>	Name	Instance ID	Instance Type	Availability Zone	Instance State	Status Checks	Alarm Status	Public DNS (IPv4)	
<input type="checkbox"/>	Blog	i-04c4fcc4f473a1d6	t2.micro	us-east-1b	running	2/2 checks ...	None	ec2-52-91-104-140.co...	
<input type="checkbox"/>	Main	i-0d90c9c4fe3dfbb8	t2.micro	us-east-1b	running	2/2 checks ...	None	ec2-54-147-242-10.co...	

2. Install Apache server on instances

Now SSH into the first instance (with name Main) and run the following commands to install and start the apache server.

```
sudo yum update -y
sudo yum install -y httpd
sudo service httpd start
sudo chkconfig httpd on
cd /var/www/html
sudo su
echo "This is the Main Website" > index.html
```

Now paste the IP address of the instance in the browser and hit Enter.


You should see something as shown in the picture below.



Now SSH into the second instance (with name Blog) and run the same commands except the last command. Instead, run the following command.

```
echo "This is the Blog Website" > blog
```

Paste the IP address of this instance with **/blog** as the suffix in the browser and hit Enter. You should see something like below.



3. Create Target Groups

A target group allows you to tell the load balancer which protocol and port will receive the traffic on the registered instances.

1. In the left navigation bar, scroll down and click on **Target Groups**.

Now click on **Create target group** at the top.

Create target group

×

Your load balancer routes requests to the targets in a target group using the target group settings that you specify, and performs health checks on the targets using the health check settings that you specify.

Target group name ⓘ

Main

Target type

☒ Instance
 ☐ IP

Protocol ⓘ

HTTP

Port ⓘ

80

VPC ⓘ

vpc-172.31.0.0/16 (My Default VPC)

Health check settings

Protocol ⓘ

HTTP

Path ⓘ

/

▶ Advanced health check settings

Cancel

Create

2. Give your target group a name **Main** and click **Create** button.

Name	Port	Protocol	Target type	Load Balancer	VPC ID	Monitoring
Main	80	HTTP	instance		vpc-	

Target group: Main

Description

Targets

Health checks

Monitoring

Tags

Edit

Registered targets

Instance ID	Name	Port	Availability Zone	Status
There are no targets registered to this target group				

Availability Zones

Availability Zone	Target count	Healthy?
There are no targets registered to this target group		

3. Now , navigate to the **Targets** tab at the bottom, click on **Edit**,select the Main instance, click **Add to registered**and click**Save**.

Register and deregister targets

Registered targets

To deregister instances, select one or more registered instances and then click Remove.

Remove

<input type="checkbox"/>	Instance	Name	Port	State	Security groups	Zone
<input type="checkbox"/>	i-0d90c9c4fe3dffbb8	Main	80	running	MyWebDMZ	us-east-1b

Instances

To register additional instances, select one or more running instances, specify a port, and then click Add. The default port is the port specified for the target group. If the instance is already registered on the specified port, you must specify a different port.

Add to registered on port 80

Search Instances

<input type="checkbox"/>	Instance	Name	State	Security groups	Zone	Subnet ID	Subnet CIDR
<input type="checkbox"/>	i-04c4fcc4f4f73a1d6	Blog	running	MyWebDMZ	us-east-1b	subnet-263ffa09	172.31.80.0/20
<input checked="" type="checkbox"/>	i-0d90c9c4fe3dffbb8	Main	running	MyWebDMZ	us-east-1b	subnet-263ffa09	172.31.80.0/20

Cancel Save

Create another target group with the name Blog and add the Blog instance to it as we did above.

4. Creating and configuring the Application Load Balancer

Now, in the left navigation scroll down and click on **Load Balancers**. Click on the **Create Load Balancer** button at the top.

ELASTIC BLOCK STORE

Volumes

Snapshots

Lifecycle Manager

NETWORK & SECURITY

Security Groups

Elastic IPs

Placement Groups

Key Pairs

Network Interfaces

LOAD BALANCING

Load Balancers

Target Groups

Create Load Balancer Actions

Filter by tags and attributes or search by keyword

<input type="checkbox"/>	Name	DNS name	State	VPC ID	Availability Zones	Type
You do not have any load balancers in this region.						

1. Choose the Application Load Balancer.

Select load balancer type

Elastic Load Balancing supports three types of load balancers: Application Load Balancers, Network Load Balancers (new), and Classic Load Balancers. Choose the load balancer type that meets your needs. [Learn more about which load balancer is right for you](#)

Application Load Balancer

HTTP HTTPS

Create

Choose an Application Load Balancer when you need a flexible feature set for your web applications with HTTP and HTTPS traffic. Operating at the request level, Application Load Balancers provide advanced routing and visibility features targeted at application architectures, including microservices and containers.

Learn more >

Network Load Balancer

TCP

Create

Choose a Network Load Balancer when you need ultra-high performance and static IP addresses for your application. Operating at the connection level, Network Load Balancers are capable of handling millions of requests per second while maintaining ultra-low latencies.

Learn more >

Classic Load Balancer

PREVIOUS GENERATION
for HTTP, HTTPS, and TCP

Create

Choose a Classic Load Balancer when you have an existing application running in the EC2-Classical network.

Learn more >

2. Give a name to your load balancer and select at least two availability zones for high availability and click on the **Next: Configure Security Settings** button.

1. Configure Load Balancer 2. Configure Security Settings 3. Configure Security Groups 4. Configure Routing 5. Register Targets 6. Review

Step 1: Configure Load Balancer

Name ⓘ my-load-balancer

Scheme ⓘ ☒ internet-facing
☐ internal

IP address type ⓘ ipv4

Listeners

A listener is a process that checks for connection requests, using the protocol and port that you configured.

Load Balancer Protocol	Load Balancer Port
HTTP	80

Add listener

Availability Zones

Specify the Availability Zones to enable for your load balancer. The load balancer routes traffic to the targets in these Availability Zones only. You can specify only one subnet per Availability Zone. You must specify subnets from at least two Availability Zones to increase the availability of your load balancer.

VPC ⓘ vpc: (172.31.0.0/16) (default)

Availability Zone	Subnet ID	Subnet IPv4 CIDR	Name
<input checked="" type="checkbox"/> us-east-1a	subnet:	172.31.0.0/20	

Cancel Next: Configure Security Settings

3. You may see a warning message but that is because we are only listening for HTTP traffic which is fine for our case, so click on the **Next: Configure Security Groups** button again.

4. Here select the existing group option and select the same security group that you assigned to the instances you launched. Once done click on **Next: Configure Routing** button.

5. In Target groups, select the existing target group. In the name select **Main** and click **Next**.

1. Configure Load Balancer 2. Configure Security Settings 3. Configure Security Groups 4. Configure Routing 5. Register Targets 6. Review

Step 4: Configure Routing

Your load balancer routes requests to the targets in this target group using the protocol and port that you specify, and performs health checks on the targets using these health check settings. Note that each target group can be associated with only one load balancer.

Target group

Target group ⓘ Existing target group

Name ⓘ Main

Target type ☒ Instance
☐ IP

Protocol ⓘ HTTP

Port ⓘ 80

Health checks

Protocol ⓘ HTTP

Path ⓘ /

► Advanced health check settings

Cancel Previous Next: Register Targets

6. Click Next again, review and click **Create**.

Name	DNS name	State	VPC ID	Availability Zones	Type
my-load-balancer	my-load-balancer-21507358...	provisioning	vpc-	us-east-1f, us-east-1b, ...	application

Load balancer: **my-load-balancer**

Description Listeners Monitoring Tags

Basic Configuration

Name: my-load-balancer ↻	Creation time: November 19, 2018 at 11:07:43 PM UTC+5:30
ARN: arn:aws:elasticloadbalancing:us-east-1: 1; b2	Hosted zone: Z35SXD0TRQ7X7K
DNS name: my-load-balancer-215073587.us-east-1.elb.amazonaws.com ⓘ (A Record)	State: provisioning
Scheme: internet-facing	VPC: vpc-
Type: application	IP address type: ipv4
	AWS WAF Web ACL:

Congratulations, you have just created an Application Load Balancer!

But we still have to configure our Blog instance so let's continue. Take a note of the **DNS name** of the Load balancer here. We will need it at the end.

7. Select the **Listeners** tab and Click on **View/edit rules**.

Name	DNS name	State	VPC ID	Availability Zones	Type
my-load-balancer	my-load-balancer-21507358...	active	vpc-	us-east-1f, us-east-1b, ...	application

Load balancer: **my-load-balancer**

Description **Listeners** Monitoring Tags

A listener checks for connection requests using its configured protocol and port, and the load balancer uses the listener rules to route requests to targets. You can add, remove, or update listeners and listener rules.

[Add listener](#) [Edit](#) [Delete](#)

Listener ID	Security policy	SSL Certificate	Rules
<input type="checkbox"/> HTTP : 80 arn:	N/A	N/A	Default: forwarding to First-Target View/edit rules

8. Click the + sign at the top to add a rule. In Add Condition select **"Path is"** and type **/blog**.

9. Then in Add Action select **Forward to** and select **Blog** and then click **Save**.

Rules

+

↑↓

−

my-load-balancer | HTTP:80

↺

?

Click a location for your new rule. Each rule must include one action of type forward, redirect, fixed response.

CancelSave

my-load-balancer | HTTP:80 (2 rules)

↑

Insert Rule

↓

RULE ID	IF (all match)	THEN
1	<div>A rule ID (ARN) is generated when you save your rule.</div> <div>✎ Path is /blog</div> <div>✎ + Add condition</div>	<div>✎ 1. Forward to Blog</div> <div>✎ + Add action</div>
last	<div>HTTP 80: default action</div> <div>This rule cannot be moved or deleted</div> <div>✓ Requests otherwise not routed</div>	<div>THEN</div> <div>Forward to Main</div>

Now, we can use the **DNS name** of our load balancer to visit the two different paths and see the results.

For /

← → ↺

ⓘ Not Secure

my-load-balancer-295044567.us-east-1.elb.amazonaws.com

This is the main website

For /blog

← → ↺

ⓘ Not Secure

my-load-balancer-295044567.us-east-1.elb.amazonaws.com/blog

This is the Blog Website